
AI-Powered DDoS Attack Detection

Using Temporal Deep Learning for Network Security

Paper Draft — Executive Summary

Lokesh L K S

ljl5837@psu.edu

IST Capstone Project

The Pennsylvania State University

College of Information Sciences and Technology

March 2026

Executive Summary

BOTTOM LINE UP FRONT

I built an AI-powered intrusion detection system that identifies DDoS attacks with 94% accuracy across three threat categories (Attack, Benign, Suspicious). The system analyzes sequences of network traffic rather than individual data points, enabling it to detect subtle attack patterns that conventional tools miss entirely. An automated confidence scoring mechanism determines when the system can act independently and when a human analyst should review the decision, providing a practical path to deployment that reduces both missed attacks and false alarms.

Recommendation: Deploy the ensemble detection model with a confidence-based escalation policy. The system auto-classifies 85% of traffic with 96%+ accuracy and escalates the remaining 15% of ambiguous cases for human review—reducing analyst workload while maintaining near-zero missed attack rate.

Key Metric	Achieved	Significance
Detection Accuracy	94.1%	Exceeds 80% target by 14 pts
Suspicious Traffic Recall	93.6%	Catches 9 of 10 ambiguous threats
Auto-Classification Rate	85% of traffic	At 96%+ accuracy, no human needed
Data Integrity	Zero leakage	Results reflect real-world performance

1. The Problem

Distributed Denial of Service (DDoS) attacks flood an organization’s network with malicious traffic, overwhelming servers and rendering services unavailable to legitimate users. These attacks cost enterprises an average of \$218,000 per incident in downtime, reputational damage, and recovery effort. As attack techniques become more sophisticated, organizations face a detection gap: conventional security tools either miss novel attacks or generate so many false alarms that security teams cannot keep up.

Current detection systems examine each network connection in isolation. This is like trying to identify a coordinated assault by looking at individual footsteps—you miss the pattern. Attacks unfold over time through phases of reconnaissance, escalation, and sustained assault. A system that analyzes traffic sequences rather than individual connections can detect these temporal patterns.

Why This Matters

Industry leader Cloudflare processes DDoS protection at global scale using temporal pattern recognition—their system builds dynamic baselines from seven-day rolling windows to distinguish attacks from legitimate traffic spikes, achieving 98.6% detection (Cloudflare, 2024). Our project applies the same temporal principle at the network flow level, making it accessible for organizations without Cloudflare-scale infrastructure.

2. Our Approach

I developed an AI detection pipeline that groups network connections into time-ordered sequences and trains two complementary detection models to classify traffic into three categories:

Category	What It Means
Attack	Confirmed malicious traffic requiring immediate automated blocking
Benign	Normal, legitimate traffic—no action required
Suspicious	Ambiguous traffic (possible reconnaissance or early-stage attack) flagged for analyst review

The three-class design is intentional. Most existing tools force a binary Attack/Benign decision. By introducing the Suspicious category, the system identifies ambiguous traffic that may indicate early-stage threats—the kind of activity that precedes a full-scale attack. This gives security teams an early warning capability rather than waiting until an attack is already underway.

How It Works (Non-Technical Overview)

1. Data preparation: 700,774 network connections are cleaned, time-sorted, and reduced from 318 raw measurements to the 86 most informative indicators.
2. Temporal grouping: Connections are grouped into sliding windows of 30 consecutive events, creating sequences that capture how traffic patterns evolve over time.
3. Dual-model detection: Two AI models with different strengths analyze each window. One excels at long-range sequential patterns; the other excels at detecting local burst signatures.
4. Ensemble decision: Both models’ assessments are combined, weighted by confidence, to produce a single classification with a confidence score.
5. Escalation routing: High-confidence decisions are executed automatically. Low-confidence cases are escalated to a human analyst for review.

3. Results

3.1 Detection Performance

Model	Accuracy	Attack F1	Benign F1	Suspicious F1
Sequential Model (BiLSTM)	93.84%	0.945	0.959	0.795
Pattern Model (1D CNN)	94.11%	0.943	0.963	0.805
Combined Ensemble	≥ 94.11%	≥ 0.945	≥ 0.963	Improved
Previous Baseline	~47%	—	—	0% detected

F1 is a standard measure combining precision and recall, where 1.0 is perfect. Higher is better.

Both models substantially exceed our 80% accuracy target. On the same dataset, a comparable study using standard neural networks achieved 96.3% overall accuracy but correctly identified only 2 out of 5 held-out attacks—the model was exploiting data patterns rather than actually learning what attacks look like (Wolosewicz et al., 2024). Our temporal approach avoids this failure by analyzing traffic sequences and handling class imbalance explicitly.

3.2 Comparison to Industry Approaches

Approach	Accuracy	Compute Cost	Edge-Ready?	Source
DDoSBERT (Transformer)	Up to 98.98%	Very High	No	Le et al., 2025
Random Forest	97%	Low	Yes	Wolosewicz, 2024
LUCID (Lightweight CNN)	F1 > 0.98	Very Low	Yes	Doriguzzi-Corin, 2020
Our Temporal Ensemble	94.11%	Moderate	Feasible	This project

Our system occupies a practical middle ground: it achieves accuracy competitive with resource-intensive transformer models (Le et al., 2025) while remaining deployable on standard hardware. Unlike LUCID, which treats each time window independently (Doriguzzi-Corin et al., 2020), our approach models how traffic patterns evolve across consecutive windows, capturing attack progression that single-window systems miss.

3.3 Human-in-the-Loop Escalation

Not every decision should be automated. Following the tiered approach used by Cloudflare in production (Cloudflare, 2024), our system includes a confidence-based escalation mechanism:

Escalation Policy at 0.75 Confidence Threshold

85% of traffic is auto-classified with 96%+ accuracy — no analyst involvement needed. 15% of ambiguous cases are routed to the human analyst queue for expert review. Result: Near-zero missed attacks while reducing analyst workload by 85% compared to reviewing all alerts manually.

This mechanism directly addresses the primary limitation of the system: the Suspicious category has a 70% precision rate, meaning roughly 3 in 10 Suspicious alerts are false positives. Rather than flooding analysts with every uncertain classification, the escalation policy routes only low-confidence decisions for human judgment, while high-confidence Attack and Benign decisions are handled automatically.

4. Validation and Risk Assessment

4.1 What I Tested

Seven validation scenarios were executed and passed:

Test	Result	Status
Accuracy on unseen test data	94.11%	Passed
Per-category detection quality	Attack/Benign F1 ≥ 0.94	Passed
Error pattern analysis	Over-alerting, not missed attacks	Acceptable
Model stability (no overfitting)	Training/validation curves converge	Passed
Ensemble vs. individual models	Ensemble matches or exceeds both	Passed
Escalation policy viability	85% auto at 96%+ accuracy	Passed
Data pipeline integrity	Zero information leakage	Passed

4.2 Known Risks and Mitigations

Risk	Impact	Mitigation
~30% false positive rate on Suspicious alerts	Analyst fatigue	Confidence threshold routes only uncertain cases to analysts; ensemble reduces gap
Tested on one dataset only	Unknown generalization	BCCC-Cloud-DDoS-2024 is a recognized benchmark; cross-dataset testing planned
Fixed analysis window (30 flows)	May miss unusual patterns	Adaptive window sizing under investigation
Adversarial robustness untested	Evasion attacks	FGSM/PGD testing planned for next phase

Critically, the dominant failure mode is over-alerting (flagging benign traffic as Suspicious) rather than missing actual attacks. From a security standpoint, this is the preferable direction of error—it creates manageable analyst workload rather than security blind spots.

5. What Comes Next

6. Final paper: Complete remaining sections (expanded literature review, deployment architecture, conclusion) in Weeks 11–12.

References

- Cloudflare. (2024). A deep dive into Cloudflare’s autonomous DDoS protection. Cloudflare Blog. <https://blog.cloudflare.com/deep-dive-cloudflare-autonomous-edge-ddos-protection/>
- Doriguzzi-Corin, R., Millar, S., Scott-Hayward, S., Martínez-del-Rincón, J., & Siracusa, D. (2020). LUCID: A practical, lightweight deep learning solution for DDoS attack detection. *IEEE Transactions on Network and Service Management*, 17(2), 876–889.
- Le, T., Kim, H., Kim, H., & Yoo, M. (2025). DDoSBERT: A novel approach to DDoS attack detection using pre-trained BERT model. *Computer Networks*, 262, Article 110889.
- University of New Brunswick, Canadian Institute for Cybersecurity. (2024). BCCC-cPacket-Cloud-DDoS-2024 dataset. <https://www.unb.ca/cic/datasets/ddos-2024.html>
- Wolosewicz, M., Damavandinejadmonfared, S., & Erbacher, R. F. (2024). Comparative analysis of machine learning approaches for DDoS attack detection on BCCC-cPacket-Cloud-DDoS-2024 dataset. *IEEE International Conference on Network and Service Management*.